



Compiler construction

Lecture 3: LLVM language and tools

Alex Gerdes

Vårtermin 2016

Chalmers University of Technology — Gothenburg University

Introduction to LLVM

Register machines



Fast but scarce

Registers are places for data inside the CPU.

- + up to 10 times faster access than to main memory
- expensive; typically just 32 of them in a 32-bit CPU

Typically, arithmetic operations, conditional jumps, etc. operate on values stored in registers.

Most modern assembly languages use registers, which correspond closely to the machine registers.

Low Level Virtual Machine (LLVM)

- LLVM is a virtual machine
- LLVM has an unbounded number of registers
- A later step does register allocation, mapping virtual registers to real machine registers

The LLVM project



The LLVM Infrastructure

- A collection of (C++) software libraries and tools to help in building compilers, debuggers, program analysers, etc.
- Tools available on Studat Linux machines
- Can also be downloaded to your own computer, see llvm.org

History

- Started as academic project at University of Illinois in 2002
- Now a large open source project with many contributors and a growing user base

Related projects

Clang C/C++ front end; aims to replace GCC

CLI MicroSoft Common Language Interface

GHC has a LLVM backend

ACM Software System Award



LLVM was the 2012 winner of the ACM Software System Award.

Previous winners include:

- | | |
|----------|--------------------|
| • VMware | • WWW |
| • Make | • TCP/IP |
| • Java | • Postscript |
| • Spin | • T _E X |
| • Coq | • Unix |
| • Apache | • ... |

The LLVM language



Characteristic features

- Three address-code with two source registers and one destination register:
`%t2 = add i32 %t0, %t1`
- One source can be a value:
`%t5 = add i32 %t3, 42`
- Instructions are typed:
`%t8 = fadd double %t6, %t7`
`store i32 %t5, i32* %r`
- New register for each result, i.e., Static Single Assignment form

Hello world in LLVM



```
@hw = internal constant [13 x i8] c"hello world\0A\00"
declare i32 @puts(i8*)

define i32 @main () {
entry: %t1 = bitcast [13 x i8]* @hw to i8*
      %t2 = call i32 @puts(i8* %t1)
      ret i32 %t2
}
```

Comments

- The string `@hw` is a global constant (global names start with an `@`-sign); note escape sequences!
- The library function `@puts` is declared, we provide its type signature
- `@hw` is cast to type of argument to `@puts`, better (type-safe) solution later

An illegal LLVM program



```
declare void @printInt(i32 %n)
define i32 @main() {
entry: %t1 = call i32 @sum(i32 100)
      call void @printInt(i32 %t1)
      ret i32 0
}

define i32 @sum (i32 %n) {
entry: %sum = i32 0
      %i = i32 0
      br label %lab1

lab1: %i = add i32 %i, 1
      %sum = add i32 %sum, %i
      %t = icmp eq i32 %i, %n
      br i1 %t, label %end, label %lab1

end:  ret i32 %sum
}
```

An illegal LLVM program



```
declare void @printInt(i32 %n)
define i32 @main() {
entry: %t1 = call i32 @sum(i32 100)
      call void @printInt(i32 %t1)
      ret i32 0
}

define i32 @sum (i32 %n) {
entry: %sum = i32 0
      %i = i32 0
      br label %lab1

lab1: %i = add i32 %i, 1
      %sum = add i32 %sum, %i
      %t = icmp eq i32 %i, %n
      br i1 %t, label %end, label %lab1

end:  ret i32 %sum
}
```

Reasons

- Not SSA form: Two assignments to `%i` and `%sum`
- There is no `%reg = val` instruction

Corrected program



```
define i32 @sum (i32 %n) {
entry: %sum = alloca i32
      store i32 0, i32* %sum
      %i = alloca i32
      store i32 0, i32* %i
      br label %lab1

lab1: %t1 = load i32* %i
      %t2 = add i32 %t1, 1
      %t3 = load i32* %sum
      %t4 = add i32 %t2, %t3
      store i32 %t2, i32* %i
      store i32 %t4, i32* %sum
      %t5 = icmp eq i32 %t2, %n
      br i1 %t5, label %end,
          label %lab1

end:  ret i32 %t4
}
```

Corrected program



```
define i32 @sum (i32 %n) {
entry: %sum = alloca i32
      store i32 0, i32* %sum
      %i = alloca i32
      store i32 0, i32* %i
      br label %lab1

lab1: %t1 = load i32* %i
      %t2 = add i32 %t1, 1
      %t3 = load i32* %sum
      %t4 = add i32 %t2, %t3
      store i32 %t2, i32* %i
      store i32 %t4, i32* %sum
      %t5 = icmp eq i32 %t2, %n
      br i1 %t5, label %end,
          label %lab1

end:  ret i32 %t4
}
```

Comments

- `%i` and `%sum` are now pointers to memory locations
- Only one assignment to any register

Problem

This program has a lot more memory traffic!

What can LLVM's optimizer do about that?

Optimizing @sum



```
> opt -mem2reg sum.ll > sumreg.bc
> llvmdis sumreg.bc
> cat sumreg.ll

define i32 @sum(i32 %n) {
entry:
  br label %lab1
lab1:
  %i.0 = phi i32 [ 0, %entry ], [ %t2, %lab1 ]
  %sum.0 = phi i32 [ 0, %entry ], [ %t4, %lab1 ]
  %t2 = add i32 %i.0, 1
  %t4 = add i32 %t2, %sum.0
  %t5 = icmp eq i32 %t2, %n
  br i1 %t5, label %end, label %lab1
end:
  ret i32 %t4
}
```

Single Static Assignment (SSA) form

- Only one assignment in the program text to each variable
- But dynamically, this assignment can be executed many times
- Many stores to a memory location are allowed
- Also, Φ (phi) instructions can be used, in the beginning of a basic block
 - Value is one of the arguments, depending on from which block control came to this block
 - Register allocation tries to keep these variables in same real register

Why SSA form?

Many code optimizations can be done more efficiently (later).

Many optimization passes

The LLVM optimizer `opt` implements many code analysis and improvement methods.

To get a default selection, give command line argument:

```
-std-compile-opts
```

Result after `opt -std-compile-opts (1/2)`

```
declare void @printInt(i32)

define i32 @main() {
entry:
    tail call void @printInt(i32 5050)
    ret i32 0
}
```

Result after `opt -std-compile-opts (2/2)`

```
define i32 @sum(i32 %n) nounwind readonly {
entry:
    %0 = shl i32 %n, 1
    %1 = add i32 %n, -1
    %2 = zext i32 %1 to i33
    %3 = add i32 %n, -2
    %4 = zext i32 %3 to i33
    %5 = mul i33 %2, %4
    %6 = lshr i33 %5, 1
    %7 = trunc i33 %6 to i32
    %8 = add i32 %0, %7
    %9 = add i32 %8, -1
    ret i32 %9
}
```

Observations

- Previous loop with execution time $O(n)$ has been optimized to code without loop, running in constant time
- Recall $1 + 2 + \dots + n = \frac{n(n+1)}{2}$, check that optimized code computes this
- Why extensions/truncations to and from 33 bits?
- What happens when n is negative?

Optimization

- `opt -std-compile-opts` includes many optimization passes
- Use `-time-passes` for an overview
- We will discuss some of these algorithms later

Part of `runtime.ll`

```
@dnl = internal constant [4 x i8] c"%d\0A\00"

declare i32 @printf(i8*, ...)

define void @printInt(i32 %x) {
entry: %t0 = getelementptr [4 x i8]* @dnl, i32 0, i32 0
    call i32 @printf(i8* %t0, i32 %x)
    ret void
}
```

We provide this file on the course web site; you just have to make sure that it is available for linking.

Linking is done by `llvm-link`

```
> llvm-link sumopt.bc runtime.bc -o a.out.bc
> llc --filetype=obj a.out.bc
> gcc a.out.o
> ./a.out
5050
```

When creating an executable file:

- Link the bitcode files with `llvm-link`.
- Compile the bitcode file to a native object file using `llc`
- Use a C compiler to link with `libc` and produce an executable

What is in `a.out.bc`



Disassemble it!

```
> cat a.out.bc | llvm-dis -  
  
; ModuleID = 'a.out.bc'  
  
@dnl = internal constant [4 x i8] c"%d\0A\00"  
  
define i32 @main() {  
entry:  
    %t0 = getelementptr [4 x i8]* @dnl, i32 0, i32 0  
    call i32 @printf(i8*, ...)* @printf(i8* %t0, i32 5050)  
    ret i32 0  
}
```

`declare i32 @printf(i8*, ...)`

¹Result slightly edited

LLVM language and tools

Types in LLVM



An incomplete list

Below t and t_i are types and n an integer literal.

- n bit integers: `in`
- `float` and `double`
- Labels: `label`
- The void type: `void`
- Functions: $t(t_1, t_2, \dots, t_n)$
- Pointer types: t^*
- Structures: $\{t_1, t_2, \dots, t_n\}$
- Arrays: $[n \times t]$

Named types and type equality



Named types

We can give names to types, for example:

```
%length = type i32  
%list = type %Node*  
%Node = type { i32, %Node* }  
  
%tree = type %Node2*  
%Node2 = type { %tree, i32, %tree }  
  
%matrix = type [ 100 x [ 100 x double ] ]
```

Type equality

LLVM uses structural equality for types.

When disassembling bitcode files that contain several structurally equal types with different names, this may give confusing results.

Identifiers



Local identifiers

Registers and named types have local names and start with a `%`-sign.

Global identifiers

Functions and global variables have global names and start with an `@`-sign.

JVALETTE does not have global variables, but you will need to define global names for string literals, as in

```
@hw = internal constant [13 x i8] c"hello world\0A\00"
```

After this definition, `@hw` has type `[13 x i8]*`.

Constants



Literals

- Integer and floating-point literals are as expected
- `true` and `false` are literals of type `i1`
- `null` is a literal of any pointer type

Aggregates

Constant expressions of structure and array types can be formed; not needed by JVALETTE.

Function definitions



Function definition form

```
define t @name(t1 x1, t2 x2, ..., tn xn) {  
  l1: block1  
  l2: block2  
  ...  
  lm: blockm  
}
```

where @name is a global name (the name of the function), the x_i are local names (the parameters) and the block_i are labeled basic blocks.

Basic blocks

A basic block is a label (l_i) followed by a colon and a sequence of LLVM instructions, each on a separate line. The last instruction must be a terminator instruction.

Function declarations



Type-checking

The LLVM assembler does type-checking. Hence it must know the types of all external functions, i.e., functions used but not defined in the compiled unit.

Simple function declaration

The basic form is: `declare t @name(t1, t2, ..., tn)`

For JAVALETTE, this is necessary for IO functions. The compiler would typically insert in each file:

```
declare void @printInt(i32)  
declare void @printDouble(double)  
declare void @printString(i8*)  
declare i32 @readInt()  
declare double @readDouble()
```

LLVM tools



llvm-as An assembler that translates llvm code to bitcode (prog.ll to prog.bc)
llvm-dis A disassembler that translates in the opposite direction
lli An interpreter/JIT compiler that executes a bitcode file containing a @main function
llvm-link A linker that links together several bitcode files
llc A compiler that translates a bitcode to native assembler or object files
opt An optimizer that optimizes bitcode; many options to decide on which optimizations to run; use -std-compile-opts to get a default selection
clang Drop-in replacement for GCC

Use of LLVM in your compiler



Default mode

Your code generator produces an assembler file (.ll). Then your main program uses system calls to first assemble this with llvm-as, optimize with opt and then link together with runtime.bc.

Other modes

More advanced and we do not recommend these for this project.

- C++ programmers can use the LLVM libraries to build in-memory representation and then output bitcode file
- Haskell programmers can access C++ libraries via Hackage package LLVM

If you want to use non-standard libraries that you haven't written yourselves, make sure to get Alex' approval first.

LLVM instructions



Basic collection

Basic JAVALETTE will only need the following instructions:

- Terminator instructions: `ret` and `br`
- Arithmetic operations:
 - For integers `add`, `sub`, `mul`, `sdiv` and `srem`
 - For doubles `fadd`, `fsub`, `fmul` and `fdiv`
- Memory access: `alloca`, `load`, `getelementptr` and `store`
- Other: `icmp`, `fcmp` and `call`

Some of the extensions will need more instructions.

Next time

Code generation for LLVM.