



Compiler construction

Lecture 6: Code generation for x86

Alex Gerdes

Vårtermin 2016

Chalmers University of Technology — Gothenburg University

This lecture



- x86 architecture
- Calling conventions
- Some x86 instructions
- From LLVM to assembler
 - Instruction selection
 - Instruction scheduling
 - Register allocation

x86 architecture

x86: assembly for a real machine



High-level view of x86

- Not a stack machine; no direct correspondence to operand stacks
- Arithmetics, etc. is done with values in registers
- Much more limited support for function calls; you need to handle return addresses, jumps, allocation of stack frames, etc. yourself
- Your code is assembled and run; no further optimization
- CISC architecture usually has few registers; straightforward code will run slowly

x86 assembler, a first example



JAVALETTE (or C)

```
> cat ex1.jl
int f (int x, int y) {
    int z = x + y;
    return z;
}
```

This might be compiled to the assembler code to the right.

NASM assembly code

```
segment .text
global f
f:
    push dword ebp
    mov  ebp, esp
    sub  esp, 4
    mov  eax, [ebp+12]
    add  eax, [ebp+8]
    mov  [ebp-4], eax
    mov  eax, [ebp-4]
    mov  esp, ebp
    pop  ebp
    ret
```

Example explained



NASM code commented

```
segment .text                ; code area
global f                     ; f has external scope
f:                             ; entry point for f
    push dword ebp           ; save caller's fp
    mov  ebp, esp            ; set our fp
    sub  esp, 4              ; allocate space for z
    mov  eax, [ebp+12]        ; move y to eax
    add  eax, [ebp+8]         ; add x to eax
    mov  [ebp-4], eax         ; move eax to z
    mov  eax, [ebp-4]         ; return value to eax
    mov  esp, ebp            ; restore caller's sp
    pop  ebp                 ; restore caller's fp
    ret                      ; pop return addr, jump
```

Intel x86 architectures



Long history

- 8086** 1978. First IBM PCs, 16 bit registers, real mode
- 80286** 1982. AT, Windows, protected mode
- 80386** 1985. 32 bit registers, virtual memory
- 80486** (Pentium, Pentium II, III, IV) 1989 – 2003. Math coprocessor, pipelining, caches, SSE, ...

Intel Core 2 2006. Multi-core

Core i3/i5/i7 2009 –

Backwards compatibility important; leading to a large set of opcodes.

Not only Intel offer x86 processors, also AMD is in the market.

Which version should you target?



x86

When speaking of the x86 architecture, one generally means register/instruction set for the 80386 (with floating-point operations).

You can compile code which would run on a 386 – or you may use SSE2 operations for a more recent version.

x86 registers



General purpose registers (32-bits)

EAX, EBX, ECX, EDX, EBP, ESP, ESI, EDI.

Conventional use:

EBP and ESP for frame pointer and stack pointer.

Segment registers

Legacy from old segmented addressing architecture.

Can be ignored in JAVALETTE compilers.

Floating-point registers

Eight 80-bit registers ST0 – ST7 organised as a stack.

Flag registers

Status registers with bits for results of comparisons, etc.

We will discuss these later.

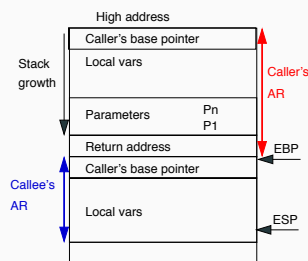
Calling convention

Data area for parameters and local variables



Runtime stack

- Contiguous memory area
- Grows from high addresses downwards
- AR layout illustrated
- EBP contains current base pointer (= frame pointer)
- ESP contains current stack pointer
- Note: We need to store return address (address of instruction to jump to on return)



Calling convention



Caller, before call

- Push params (in reverse order)
- Push return address
- Jump to callee entry

```
push dword paramn
...
push dword param1
call f
```

Calling convention



Caller, before call

- Push params (in reverse order)
 - Push return address
 - Jump to callee entry
- ```
push dword paramn
...
push dword param1
call f
```

### Callee, on entry

- Push caller's base pointer
  - Update current base pointer
  - Allocate space for locals
- ```
push dword ebp
mov ebp, esp
sub esp, localbytes
```

Calling convention



Caller, before call

- Push params (in reverse order)
 - Push return address
 - Jump to callee entry
- ```
push dword paramn
...
push dword param1
call f
```

### Callee, on entry

- Push caller's base pointer
  - Update current base pointer
  - Allocate space for locals
- ```
push dword ebp
mov ebp, esp
sub esp, localbytes
```

Callee, on exit

- Restore base and stack ptr
 - Pop return address and jump
- ```
mov esp, ebp
pop ebp
ret
```

## Calling convention



### Caller, before call

- Push params (in reverse order)
  - Push return address
  - Jump to callee entry
- ```
push dword paramn
...
push dword param1
call f
```

Callee, on entry

- Push caller's base pointer
 - Update current base pointer
 - Allocate space for locals
- ```
push dword ebp
mov ebp, esp
sub esp, localbytes
```

### Callee, on exit

- Restore base and stack ptr
  - Pop return address and jump
- ```
mov esp, ebp
pop ebp
ret
```

Caller, after call

- Pop parameters
- ```
add esp, parambytes
```

## Calling convention



### Caller, before call

- Push params (in reverse order)
  - Push return address
  - Jump to callee entry
- ```
push dword paramn
...
push dword param1
call f
```

Callee, on entry

- Push caller's base pointer
 - Update current base pointer
 - Allocate space for locals
- ```
enter localbytes, 0
```

### Callee, on exit

- Restore base and stack ptr
  - Pop return address and jump
- ```
leave
ret
```

Caller, after call

- Pop parameters
- ```
add esp, parambytes
```

## Parameters, local variables and return values



### Parameters

- In the callee code, integer parameter 1 has address EBP+8, parameter 2 EBP+12, etc.
- Parameter values accessed with indirect addressing: [EBP+8], etc.
- Double parameters require 8 bytes
- Here EBP+n means "(address stored in EBP) + n"

### Local variables

- First local var is at address EBP-4, etc.
- Local vars are conventionally addressed relative to EBP, not ESP
- Again, refer to vars by indirect addressing: [EBP-4], etc.

### Return values

Integer and boolean values are returned in EAX, doubles in ST0

## Register usage



### Scratch registers (caller save)

EAX, ECX and EDX must be saved by caller before call, if used; can be freely used by callee.

### Callee save register

EBX, ESI, EDI, EBP, ESP.

For EBP and ESP, this is handled in the code patterns.

### Note

- What we have described is one common calling convention for 32-bit x86, called cdecl
- Other conventions exist, but we omit them

## Assemblers for x86



### Several alternatives

- Several assemblers for x86 exist, with different syntax
- We will use NASM, the Netwide Assembler, which is available for several platforms
- We also recommend Paul Carter's book and examples; follow link from course website
- Some syntax differences to the GNU assembler:
  - GNU uses `%eax` etc. as register names
  - For two-argument instructions, the operands have opposite order!
  - Different syntax for indirect addressing

If you use `gcc -S ex.c`, you will get GNU syntax

## Example: GNU syntax



### First example, revisited

```
> gcc -c ex1.c
> objdump -d ex1.o
```

```
ex1.o: file format elf32-i386
Disassembly of section .text:
```

```
00000000 <f>:
0: 55 push %ebp
1: 89 e5 mov %esp,%ebp
3: 8b 45 0c mov 0xc(%ebp),%eax
6: 03 45 08 add 0x8(%ebp),%eax
9: c9 leave %eax
a: c3 ret
```

## Assembler

## Integer arithmetic; two-address code



### Addition, subtraction and multiplication

```
add dest, src ; dest := dest + src
sub dest, src ; dest := dest - src
imul dest, src ; dest := dest * src
```

Operands can be values in registers or in memory; `src` also a literal.

### Division – one-address code

```
idiv denom
(eax, edx) := ((edx:eax) / denom, (edx:eax) % denom)
```

- The numerator is the 64-bit value EDX:EAX (no other choices)
- Both div and mod are performed; results in EAX resp. EDX
- EDX must be zeroed before division

## Example



### JAVALETTE program

```
int main () {
 printString "Input a number: ";
 int n = readInt();
 printInt(2 * n);
 return 0;
}
```

The above code could be translated as follows (slightly optimized to fit on slide).

### Code for main

```
push dword ebp
mov ebp, esp
push str1
call printString
add esp, 4
call readInt
imul eax, 2
push eax
call printInt
add esp, 4
mov eax, 0
leave
ret
```

## Example, continued



### Complete file

```
extern printString, printInt
extern readInt

segment .data
str1 db "Input a number: "

segment .text
global main

main:
; code from previous slide
```

### Comments

- IO functions are external; we will come back to that
- The `.data` segment contains constants such as `str1`
- The `.text` segment contains code
- The `global` declaration gives `main` external scope (can be called from code outside this file)

## Floating-point arithmetic in x86



### Moving numbers (selection)

`fld src` Pushes value in `src` on fp stack

`fild src` Pushes integer value in `src` on fp stack

`fstp dest` Stores top of fp stack in `dest` and pops

Both `src` and `dest` can be fp register or memory reference.

### Arithmetic (selection)

`fadd src` Adds `src` to `ST0`

`fadd to dest` Adds `ST0` to `dest`

`faddp dest` Adds `ST0` to `dest`, then pop

Similar variants for `fsub`, `fmul` and `fdiv`.

## Floating-point arithmetic in SSE2



### New registers

- 128-bit registers `XMM0–XMM7` (later also `XMM8–XMM15`)
- Each can hold two double precision floats or four single-precision floats
- SIMD operations for arithmetic

### Arithmetic instructions

- Two-address code, `ADDSD`, `MULSD`, etc.
- SSE2 fp code similar to integer arithmetic

## Control flow



### Integer comparisons

- `cmp v1 v2`
- `v1 - v2` is computed and bits in the flag register are set:
  - ZF is set iff value is zero
  - OF is set iff result overflows
  - SF is set iff result is negative

## Control flow



### Integer comparisons

- `cmp v1 v2`
- `v1 - v2` is computed and bits in the flag register are set:
  - ZF is set iff value is zero
  - OF is set iff result overflows
  - SF is set iff result is negative

### Branch instructions (selection)

- `JZ lab` branches if ZF is set
- `JL lab` branches if SF is set
- Similarly for the other relations between `v1` and `v2`
- `fcomi src` compares `ST0` and `src` and sets flags; can be followed by branching as above

## One more example



### JAVALETTE (or C)

```
int sum(int n) {
 int res = 0;
 int i = 0;
 while (i < n) {
 res = res + i;
 i++;
 }
 return res;
}
```

### Naive assembler

```
sum: enter 8, 0
 mov [ebp-4], 0
 mov [ebp-8], 0
 jmp L2
L3: mov eax, [ebp-8]
 add [ebp-4], eax
 inc [ebp-8]
L2: mov eax, [ebp-8]
 cmp eax, [ebp+8]
 jl L3
 mov eax, [ebp-4]
 leave
 ret
```

## How to do an x86 backend



### Starting point

Two alternatives:

- From LLVM code (requires your basic backend to generate LLVM code as a data structure, not directly as strings); will generate many local vars
- From AST's generated by the frontend (means a lot of code common with LLVM backend)

### Variables

In either case, your code will contain a lot of variables/virtual registers. Possible approaches:

- Treat these as local vars, storing to and fetching from stack at each access; gives really slow code
- Do register allocation<sup>1</sup>; much better code

<sup>1</sup>Future lecture

## Input and output



### A simple proposal

Define `printInt`, `readInt`, etc. in C. Then link this file together with your object files using `gcc`.

Alternative: Compile `runtime.ll` with `llvm-as` and `llc` to get `runtime.s`; this can be given to `gcc` as below.

### Linux building

To assemble a NASM file to `file.o`:

```
nasm -f elf file.asm
```

To link:

```
gcc file.o runtime.s
```

Result is executable `a.out`

### More info

Paul Carter's book (link on course web site) gives more info.

## From LLVM to assembler

## From LLVM to assembler



### Several stages

- Instruction selection
- Instruction scheduling
- SSA-based optimizations
- Register allocation
- Prolog/epilog code (AR management)
- Code emission

### Target-independent generation

Also much of this is done in target-independent ways and using general algorithms operating on target descriptions.

## Native code generation, revisited



### More complications

So far, we have ignored some important concerns in code generation:

- The instruction set in real-world processors typically offer many different ways to achieve the same effect. Thus, when translating an IR program to native code we must do instruction selection, i.e., choose between available alternatives.
- Often an instruction sequence contain independent parts that can be executed in arbitrary order. Different orders may take very different time; thus a code generator should do instruction scheduling.

Both these task are complex and interact with register allocation.

In LLVM, these tasks are done by the native code generator `llc` and the JIT compiler in `lli`.

## Instruction selection



### Further observations

- Instruction selection for RISC machines generally simpler than for CISC machines
- The number of translation possibilities grow (combinatorially) as one considers larger chunks of IR code for translation

### Pattern matching

The IR code can be seen as a pattern matching problem: The native instructions are seen as patterns; instruction selection is the problem to cover the IR code by patterns.

## Instruction selection



### Further observations

- Instruction selection for RISC machines generally simpler than for CISC machines
- The number of translation possibilities grow (combinatorially) as one considers larger chunks of IR code for translation

### Pattern matching

The IR code can be seen as a pattern matching problem: The native instructions are seen as patterns; instruction selection is the problem to cover the IR code by patterns.


### Two approaches

- Tree pattern matching: think of IR code as tree
- Peephole matching: think of IR code as sequence

## Tree pattern matching, an example



$a[i] := x$  as tree IR code




- $a$  and  $x$  local vars,  $i$  in register
- $a$  is pointer to first element

## Tree pattern matching, an example



$a[i] := x$  as tree IR code



### Algorithm outline

- Represent native instructions as patterns, or tree fragments
- Tile the IR tree using these patterns so that all nodes in the tree are covered
- Output the sequence of instructions corresponding to the tiling

- $a$  and  $x$  local vars,  $i$  in register
- $a$  is pointer to first element

## A simple instruction set

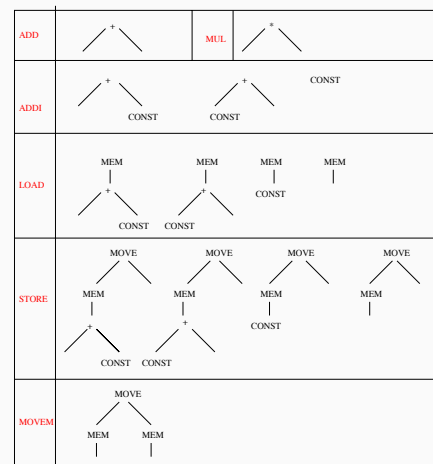


|       |                             |
|-------|-----------------------------|
| ADD   | $r_i \leftarrow r_j + r_k$  |
| MUL   | $r_i \leftarrow r_j * r_k$  |
| SUB   | $r_i \leftarrow r_j - r_k$  |
| DIV   | $r_i \leftarrow r_j / r_k$  |
| ADDI  | $r_i \leftarrow r_j + c$    |
| SUBI  | $r_i \leftarrow r_j - c$    |
| LOAD  | $r_i \leftarrow M[r_j + c]$ |
| STORE | $M[r_j + c] \leftarrow r_i$ |
| MOVEM | $M[r_j] \leftarrow M[r_i]$  |

### Notes

- We consider only arithmetic and memory instructions (no jumps!)

## Identifying patterns (incomplete)



## Instruction scheduling, background



### Simple-minded, old-fashioned view of processor

Fetch an instruction, decode it, fetch operands, perform operation, store result. Then fetch next operation, ...

### Modern processors

- Several instructions under execution concurrently
- Memory system cause delays, with operations waiting for data
- Similar problems for results from arithmetic operations, that may take several cycles

### Consequence

Important to understand data dependencies and order instructions advantageously.

## Instruction scheduling, example



### Example (from Cooper)

$w = w * 2 * x * y * z$

- Memory op takes 3 cycles, mult 2 cycles, add one cycle
- One instruction can be issued each cycle, if data available

## Instruction scheduling, example



### Example (from Cooper)

$w = w * 2 * x * y * z$

- Memory op takes 3 cycles, mult 2 cycles, add one cycle
- One instruction can be issued each cycle, if data available

### Schedule 1

```
r1 <- M [fp + @w]
r1 <- r1 + r1
r2 <- M [fp + @x]
r1 <- r1 * r2
r2 <- M [fp + @y]
r1 <- r1 * r2
r2 <- M [fp + @z]
r1 <- r1 * r2
M [fp + @w] <- r1
```

## Instruction scheduling, example



### Example (from Cooper)

$w = w * 2 * x * y * z$

- Memory op takes 3 cycles, mult 2 cycles, add one cycle
- One instruction can be issued each cycle, if data available

### Schedule 1

```
r1 <- M [fp + @w]
r1 <- r1 + r1
r2 <- M [fp + @x]
r1 <- r1 * r2
r2 <- M [fp + @y]
r1 <- r1 * r2
r2 <- M [fp + @z]
r1 <- r1 * r2
M [fp + @w] <- r1
```

### Schedule 2

```
r1 <- M [fp + @w]
r2 <- M [fp + @x]
r3 <- M [fp + @y]
r1 <- r1 + r1
r2 <- r1 * r2
r2 <- M [fp + @z]
r1 <- r1 * r3
r1 <- r1 * r2
M [fp + @w] <- r1
```

## Instruction scheduling



### Comments

- Problem is NP-complete for realistic architectures
- Common technique is list scheduling: greedy algorithm for scheduling a basic block
- Builds graph describing data dependencies between instructions and schedules instructions from ready list of instructions with available operands

### Interaction

Despite interaction between selection, scheduling and register allocation, these are typically handled independently (and in this order).

## x86 backend extension



### Comments

- Increased to two credits
- Acts as a 'multiplier' for other extensions